

```

from __future__ import print_function

from keras.models import Model
from keras.layers import Input, LSTM, Dense, Flatten, Attention
import numpy as np
import pandas as pd
pd.set_option('display.max_columns', None)

batch_size = 64 # Batch size for training.
epochs = 50 # Number of epochs to train for.
latent_dim = 256 # Latent dimensionality of the encoding space.
num_samples = 3000 # Number of samples to train on.
# Path to the data txt file on disk.
data_path = 'swe.txt'

```

```

##Question 1
# Vectorize the data.
input_texts = []
target_texts = []
input_characters = set()
target_characters = set()
with open(data_path, 'r', encoding='utf-8') as f:
    lines = f.read().split('\n')

```

```
encoder_input_data
```

```
#now need to create model for character level model
```

```

from __future__ import print_function
import os
from keras.models import Model
from keras.layers import Input, LSTM, Dense
import numpy as np
import pandas as pd
from matplotlib import pyplot
!pip install contractions
pd.set_option('display.max_columns', None)

```

Successfully installed anyascii-0.3.0 contractions-0.0.55 pyahocorasick-1.4.2 textsearch

```
import numpy as np
```

```

import typing
from typing import Any, Tuple

```

```
import tensorflow as tf
from tensorflow.keras.layers.experimental import preprocessing
!pip install tensorflow_text
import tensorflow_text as tf_text
```

```
import matplotlib.pyplot as plt
import matplotlib.ticker as ticker
```

```
Installing collected packages: tensorflow-text
Successfully installed tensorflow-text-2.6.0
```

```
##@title Shape checker
```

## Shape checker

```
class ShapeChecker():
    def __init__(self):
        # Keep a cache of every axis-name seen
        self.shapes = {}

    def __call__(self, tensor, names, broadcast=False):
        if not tf.executing_eagerly():
            return

        if isinstance(names, str):
            names = (names,)

        shape = tf.shape(tensor)
        rank = tf.rank(tensor)

        if rank != len(names):
            raise ValueError(f'Rank mismatch:\n'
                             f'    found {rank}: {shape.numpy()}\n'
                             f'    expected {len(names)}: {names}\n')

        for i, name in enumerate(names):
            if isinstance(name, int):
                old_dim = name
            else:
                old_dim = self.shapes.get(name, None)
            new_dim = shape[i]

            if (broadcast and new_dim == 1):
                continue

            if old_dim is None:
                # If the axis name is new, add its length to the cache.
                self.shapes[name] = new_dim
                continue

            if new_dim != old_dim:
                raise ValueError(f"Shape mismatch for dimension: '{name}'\n")
```

```
f"      found: {new_dim}\n"
f"      mounted: {old_dim}\n"
```

```
from google.colab import drive
drive.mount('/content/drive')
```

```
Mounted at /content/drive
```

```
use_builtins = True
```

```
import contractions
input_texts2 = []
target_texts2 = []
with open('swe.txt', 'r', encoding='utf-8') as f:
    lines = f.read().split('\n')
for line in lines[: min(num_samples, len(lines) - 1)]:
    input_text2, target_text2, useless2 = line.split('\t')
    #target_text2 = '\t' + target_text2 + '\n'
    input_texts2.append(input_text2)
    target_texts2.append(target_text2)
```

```
input_texts3= tf.convert_to_tensor(input_texts2)
target_texts3= tf.convert_to_tensor(target_texts2)
```

```
input_texts3.shape
```

```
TensorShape([3000])
```

```
BUFFER_SIZE = len(input_texts2)
BATCH_SIZE = 64
```

```
dataset = tf.data.Dataset.from_tensor_slices((input_texts3, target_texts3)).shuffle(BUFFER_SIZE)
dataset = dataset.batch(BATCH_SIZE)
```

```
for example_input_batch, example_target_batch in dataset.take(4):
    print(example_input_batch[:5])
    print()
    print(example_target_batch[:5])
    break
```

```
def tf_lower_and_split_punct(text):
    # Split accented characters.
```

```

text = tf_text.normalize_utf8(text, 'NFKD')
text = tf.strings.lower(text)
# Keep space, a to z, and select punctuation.
text = tf.strings.regex_replace(text, '[^ a-z.?!,:;]', '')
# Add spaces around punctuation.
text = tf.strings.regex_replace(text, '[.?!,:;]', r' \0 ')
# Strip whitespace.
text = tf.strings.strip(text)

text = tf.strings.join(['[START]', text, '[END]'], separator=' ')
return text

max_vocab_size = 5000

input_text_processor = preprocessing.TextVectorization(
    standardize=tf_lower_and_split_punct,
    max_tokens=max_vocab_size)

input_text_processor.adapt(input_texts3)

# Here are the first 10 words from the vocabulary:
input_text_processor.get_vocabulary()[:10]

['', '[UNK]', '[START]', '[END]', '.', 'i', '?', 'tom', 'it', 'you']

output_text_processor = preprocessing.TextVectorization(
    standardize=tf_lower_and_split_punct,
    max_tokens=max_vocab_size)

output_text_processor.adapt(target_texts3)
output_text_processor.get_vocabulary()[:10]

['', '[UNK]', '[START]', '[END]', '.', 'jag', 'ar', 'tom', '?', 'det']

#dont use this???
input_tokens = input_text_processor(input_texts3)
input_tokens[:3, :10]

<tf.Tensor: shape=(3, 8), dtype=int64, numpy=
array([[ 2, 33,  4,  3,  0,  0,  0,  0],
       [ 2, 133, 18,  3,  0,  0,  0,  0],
       [ 2, 70,  6,  3,  0,  0,  0,  0]])>

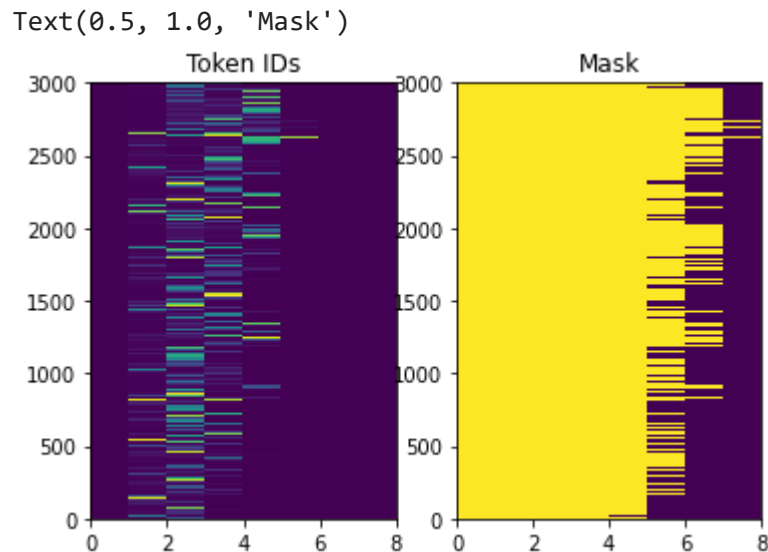
input_vocab = np.array(input_text_processor.get_vocabulary())
tokens = input_vocab[input_tokens[0].numpy()]
''.join(tokens)

'[START] go . [END]'

```

```
plt.subplot(1, 2, 1)
plt.pcolormesh(input_tokens)
plt.title('Token IDs')

plt.subplot(1, 2, 2)
plt.pcolormesh(input_tokens != 0)
plt.title('Mask')
```



```
embedding_dim = 256
units = 1024
```

```
class Encoder(tf.keras.layers.Layer):
    def __init__(self, input_vocab_size, embedding_dim, enc_units):
        super(Encoder, self).__init__()
        self.enc_units = enc_units
        self.input_vocab_size = input_vocab_size

        # The embedding layer converts tokens to vectors
        self.embedding = tf.keras.layers.Embedding(self.input_vocab_size,
                                                    embedding_dim)

        # The GRU RNN layer processes those vectors sequentially.
        self.gru = tf.keras.layers.GRU(self.enc_units,
                                         # Return the sequence and state
                                         return_sequences=True,
                                         return_state=True,
                                         recurrent_initializer='glorot_uniform')

    def call(self, tokens, state=None):
        shape_checker = ShapeChecker()
        shape_checker(tokens, ('batch', 's'))
```

```

# 2. The embedding layer looks up the embedding for each token.
vectors = self.embedding(tokens)
shape_checker(vectors, ('batch', 's', 'embed_dim'))

# 3. The GRU processes the embedding sequence.
#   output shape: (batch, s, enc_units)
#   state shape: (batch, enc_units)
output, state = self.gru(vectors, initial_state=state)
shape_checker(output, ('batch', 's', 'enc_units'))
shape_checker(state, ('batch', 'enc_units'))

# 4. Returns the new sequence and its state.
return output, state

# Convert the input text to tokens.
input_tokens2 = input_text_processor(input_texts3)

# Encode the input sequence.
encoder = Encoder(input_text_processor.vocabulary_size(),
                  embedding_dim, units)
example_enc_output, example_enc_state = encoder(input_tokens2)

print(f'Input batch tokens, shape (batch, s): {input_tokens2.shape}')
print(f'Encoder output, shape (batch, s, units): {example_enc_output.shape}')
print(f'Encoder state, shape (batch, units): {example_enc_state.shape}')

    Input batch tokens, shape (batch, s): (3000, 8)
    Encoder output, shape (batch, s, units): (3000, 8, 1024)
    Encoder state, shape (batch, units): (3000, 1024)

class BahdanauAttention(tf.keras.layers.Layer):
    def __init__(self, units):
        super().__init__()
        # For Eqn. (4), the Bahdanau attention
        self.W1 = tf.keras.layers.Dense(units, use_bias=False)
        self.W2 = tf.keras.layers.Dense(units, use_bias=False)

        self.attention = tf.keras.layers.AdditiveAttention()

    def call(self, query, value, mask):
        shape_checker = ShapeChecker()
        shape_checker(query, ('batch', 't', 'query_units'))
        shape_checker(value, ('batch', 's', 'value_units'))
        shape_checker(mask, ('batch', 's'))

        # From Eqn. (4), `W1@ht`.
        w1_query = self.W1(query)
        shape_checker(w1_query, ('batch', 't', 'attn_units'))

        # From Eqn. (4), `W2@hs`.
        w2_key = self.W2(value)

```

```

shape_checker(w2_key, ('batch', 's', 'attn_units'))

query_mask = tf.ones(tf.shape(query)[: -1], dtype=bool)
value_mask = mask

context_vector, attention_weights = self.attention(
    inputs = [w1_query, value, w2_key],
    mask=[query_mask, value_mask],
    return_attention_scores = True,
)
shape_checker(context_vector, ('batch', 't', 'value_units'))
shape_checker(attention_weights, ('batch', 't', 's'))

return context_vector, attention_weights

attention_layer = BahdanauAttention(units)

(input_tokens != 0).shape

TensorShape([3000, 8])

# Later, the decoder will generate this attention query
example_attention_query = tf.random.normal(shape=[len(input_tokens), 2, 10])

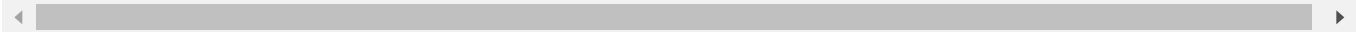
# Attend to the encoded tokens

context_vector, attention_weights = attention_layer(
    query=example_attention_query,
    value=example_enc_output,
    mask=(input_tokens != 0))

print(f'Attention result shape: (batch_size, query_seq_length, units):          {context_vec
print(f'Attention weights shape: (batch_size, query_seq_length, value_seq_length): {attention

Attention result shape: (batch_size, query_seq_length, units):          (3000, 2, 1024)
Attention weights shape: (batch_size, query_seq_length, value_seq_length): (3000, 2, 8)

```



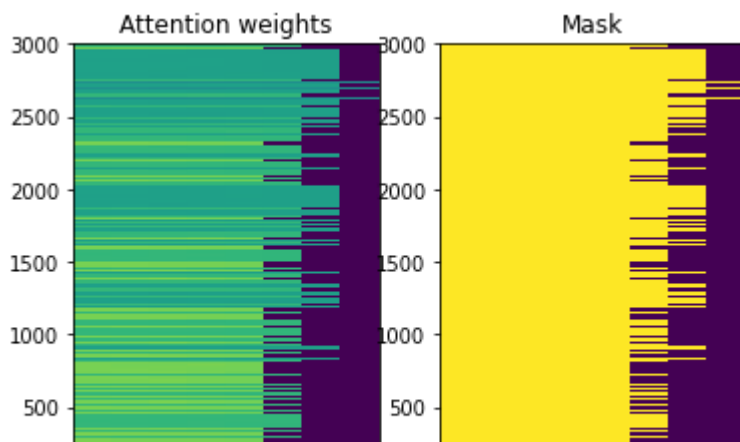
```

plt.subplot(1, 2, 1)
plt.pcolormesh(attention_weights[:, 0, :])
plt.title('Attention weights')

plt.subplot(1, 2, 2)
plt.pcolormesh(input_tokens != 0)
plt.title('Mask')

```

Text(0.5, 1.0, 'Mask')



```
attention_weights.shape
```

```
TensorShape([3000, 2, 8])
```

```
attention_slice = attention_weights[0, 0].numpy()
attention_slice = attention_slice[attention_slice != 0]
```

```
#@title
```

```
plt.suptitle('Attention weights for one sequence')
```

```
plt.figure(figsize=(12, 6))
```

```
a1 = plt.subplot(1, 2, 1)
```

```
plt.bar(range(len(attention_slice)), attention_slice)
```

```
# freeze the xlim
```

```
plt.xlim(plt.xlim())
```

```
plt.xlabel('Attention weights')
```

```
a2 = plt.subplot(1, 2, 2)
```

```
plt.bar(range(len(attention_slice)), attention_slice)
```

```
plt.xlabel('Attention weights, zoomed')
```

```
# zoom in
```

```
top = max(a1.get_ylim())
```


```
zoom = 0.85*top
```

```
a2.set_ylim([0.90*top, top])
```

```
a1.plot(a1.get_xlim(), [zoom, zoom], color='k')
```



[<matplotlib.lines.Line2D at 0x7f512b39d9d0>]  
 <Figure size 432x288 with 0 Axes>



```
class Decoder(tf.keras.layers.Layer):
    def __init__(self, output_vocab_size, embedding_dim, dec_units):
        super(Decoder, self).__init__()
        self.dec_units = dec_units
        self.output_vocab_size = output_vocab_size
        self.embedding_dim = embedding_dim

        # For Step 1. The embedding layer convets token IDs to vectors
        self.embedding = tf.keras.layers.Embedding(self.output_vocab_size,
                                                    embedding_dim)

        # For Step 2. The RNN keeps track of what's been generated so far.
        self.gru = tf.keras.layers.GRU(self.dec_units,
                                         return_sequences=True,
                                         return_state=True,
                                         recurrent_initializer='glorot_uniform')

        # For step 3. The RNN output will be the query for the attention layer.
        self.attention = BahdanauAttention(self.dec_units)

        # For step 4. Eqn. (3): converting `ct` to `at`
        self.Wc = tf.keras.layers.Dense(dec_units, activation=tf.math.tanh,
                                         use_bias=False)

        # For step 5. This fully connected layer produces the logits for each
        # output token.
        self.fc = tf.keras.layers.Dense(self.output_vocab_size)

class DecoderInput(typing.NamedTuple):
    new_tokens: Any
    enc_output: Any
    mask: Any

class DecoderOutput(typing.NamedTuple):
    logits: Any
    attention_weights: Any
```

```

def call(self,
        inputs: DecoderInput,
        state=None) -> Tuple[DecoderOutput, tf.Tensor]:
    shape_checker = ShapeChecker()
    shape_checker(inputs.new_tokens, ('batch', 't'))
    shape_checker(inputs.enc_output, ('batch', 's', 'enc_units'))
    shape_checker(inputs.mask, ('batch', 's'))

    if state is not None:
        shape_checker(state, ('batch', 'dec_units'))

    # Step 1. Lookup the embeddings
    vectors = self.embedding(inputs.new_tokens)
    shape_checker(vectors, ('batch', 't', 'embedding_dim'))

    # Step 2. Process one step with the RNN
    rnn_output, state = self.gru(vectors, initial_state=state)

    shape_checker(rnn_output, ('batch', 't', 'dec_units'))
    shape_checker(state, ('batch', 'dec_units'))

    # Step 3. Use the RNN output as the query for the attention over the
    # encoder output.
    context_vector, attention_weights = self.attention(
        query=rnn_output, value=inputs.enc_output, mask=inputs.mask)
    shape_checker(context_vector, ('batch', 't', 'dec_units'))
    shape_checker(attention_weights, ('batch', 't', 's'))

    # Step 4. Eqn. (3): Join the context_vector and rnn_output
    # [ct; ht] shape: (batch t, value_units + query_units)
    context_and_rnn_output = tf.concat([context_vector, rnn_output], axis=-1)

    # Step 4. Eqn. (3): `at = tanh(Wc@[ct; ht])`
    attention_vector = self.Wc(context_and_rnn_output)
    shape_checker(attention_vector, ('batch', 't', 'dec_units'))

    # Step 5. Generate logit predictions:
    logits = self.fc(attention_vector)
    shape_checker(logits, ('batch', 't', 'output_vocab_size'))

    return DecoderOutput(logits, attention_weights), state

```

```
Decoder.call = call
```

```

decoder = Decoder(output_text_processor.vocabulary_size(),
                  embedding_dim, units)

```

```

# Convert the target sequence, and collect the "[START]" tokens
example_output_tokens = output_text_processor(target_texts3)

```

```

start_index = output_text_processor.get_vocabulary().index('[START]')
first_token = tf.constant([[start_index]] * example_output_tokens.shape[0])

# Run the decoder
dec_result, dec_state = decoder(
    inputs = DecoderInput(new_tokens=first_token,
                           enc_output=example_enc_output,
                           mask=(input_tokens != 0)),
    state = example_enc_state
)

print(f'logits shape: (batch_size, t, output_vocab_size) {dec_result.logits.shape}')
print(f'state shape: (batch_size, dec_units) {dec_state.shape}')

logits shape: (batch_size, t, output_vocab_size) (3000, 1, 1665)
state shape: (batch_size, dec_units) (3000, 1024)

sampled_token = tf.random.categorical(dec_result.logits[:, 0, :], num_samples=1)

vocab = np.array(output_text_processor.get_vocabulary())
first_word = vocab[sampled_token.numpy()]
first_word[:5]

dec_result, dec_state = decoder(
    DecoderInput(sampled_token,
                  example_enc_output,
                  mask=(input_tokens != 0)),
    state=dec_state)

sampled_token = tf.random.categorical(dec_result.logits[:, 0, :], num_samples=1)
first_word = vocab[sampled_token.numpy()]
first_word[:5]

class MaskedLoss(tf.keras.losses.Loss):
    def __init__(self):
        self.name = 'masked_loss'
        self.loss = tf.keras.losses.SparseCategoricalCrossentropy(
            from_logits=True, reduction='none')

    def __call__(self, y_true, y_pred):
        shape_checker = ShapeChecker()
        shape_checker(y_true, ('batch', 't'))
        shape_checker(y_pred, ('batch', 't', 'logits'))

        # Calculate the loss for each item in the batch.
        loss = self.loss(y_true, y_pred)

```

```

    shape_checker(loss, ('batch', 't'))

    # Mask off the losses on padding.
    mask = tf.cast(y_true != 0, tf.float32)
    shape_checker(mask, ('batch', 't'))
    loss *= mask

    # Return the total.
    return tf.reduce_sum(loss)

class TrainTranslator(tf.keras.Model):
    def __init__(self, embedding_dim, units,
                  input_text_processor,
                  output_text_processor,
                  use_tf_function=True):
        super().__init__()
        # Build the encoder and decoder
        encoder = Encoder(input_text_processor.vocabulary_size(),
                          embedding_dim, units)
        decoder = Decoder(output_text_processor.vocabulary_size(),
                          embedding_dim, units)

        self.encoder = encoder
        self.decoder = decoder
        self.input_text_processor = input_text_processor
        self.output_text_processor = output_text_processor
        self.use_tf_function = use_tf_function
        self.shape_checker = ShapeChecker()

    def train_step(self, inputs):
        self.shape_checker = ShapeChecker()
        if self.use_tf_function:
            return self._tf_train_step(inputs)
        else:
            return self._train_step(inputs)

    def _preprocess(self, input_text, target_text):
        self.shape_checker(input_text, ('batch',))
        self.shape_checker(target_text, ('batch',))

        # Convert the text to token IDs
        input_tokens = self.input_text_processor(input_text)
        target_tokens = self.output_text_processor(target_text)
        self.shape_checker(input_tokens, ('batch', 's'))
        self.shape_checker(target_tokens, ('batch', 't'))

        # Convert IDs to masks.
        input_mask = input_tokens != 0
        self.shape_checker(input_mask, ('batch', 's'))

```

```

target_mask = target_tokens != 0
self.shape_checker(target_mask, ('batch', 't'))

return input_tokens, input_mask, target_tokens, target_mask

TrainTranslator._preprocess = _preprocess

def _train_step(self, inputs):
    input_text, target_text = inputs

    (input_tokens, input_mask,
     target_tokens, target_mask) = self._preprocess(input_text, target_text)

    max_target_length = tf.shape(target_tokens)[1]

    with tf.GradientTape() as tape:
        # Encode the input
        enc_output, enc_state = self.encoder(input_tokens)
        self.shape_checker(enc_output, ('batch', 's', 'enc_units'))
        self.shape_checker(enc_state, ('batch', 'enc_units'))

        # Initialize the decoder's state to the encoder's final state.
        # This only works if the encoder and decoder have the same number of
        # units.
        dec_state = enc_state
        loss = tf.constant(0.0)

        for t in tf.range(max_target_length-1):
            # Pass in two tokens from the target sequence:
            # 1. The current input to the decoder.
            # 2. The target for the decoder's next prediction.
            new_tokens = target_tokens[:, t:t+2]
            step_loss, dec_state = self._loop_step(new_tokens, input_mask,
                                                    enc_output, dec_state)

            loss = loss + step_loss

        # Average the loss over all non padding tokens.
        average_loss = loss / tf.reduce_sum(tf.cast(target_mask, tf.float32))

    # Apply an optimization step
    variables = self.trainable_variables
    gradients = tape.gradient(average_loss, variables)
    self.optimizer.apply_gradients(zip(gradients, variables))

    # Return a dict mapping metric names to current value
    return {'batch_loss': average_loss}

TrainTranslator._train_step = _train_step

```

```

def _loop_step(self, new_tokens, input_mask, enc_output, dec_state):
    input_token, target_token = new_tokens[:, 0:1], new_tokens[:, 1:2]

    # Run the decoder one step.
    decoder_input = DecoderInput(new_tokens=input_token,
                                  enc_output=enc_output,
                                  mask=input_mask)

    dec_result, dec_state = self.decoder(decoder_input, state=dec_state)
    self.shape_checker(dec_result.logits, ('batch', 't1', 'logits'))
    self.shape_checker(dec_result.attention_weights, ('batch', 't1', 's'))
    self.shape_checker(dec_state, ('batch', 'dec_units'))

    # `self.loss` returns the total for non-padded tokens
    y = target_token
    y_pred = dec_result.logits
    step_loss = self.loss(y, y_pred)

    return step_loss, dec_state

```

```
TrainTranslator._loop_step = _loop_step
```

```

translator = TrainTranslator(
    embedding_dim, units,
    input_text_processor=input_text_processor,
    output_text_processor=output_text_processor,
    use_tf_function=False)

```

```

# Configure the loss and optimizer
translator.compile(
    optimizer=tf.optimizers.Adam(),
    loss=MaskedLoss(),
)

```

```
np.log(output_text_processor.vocabulary_size())
```

```
7.417580402414544
```

```

%%time
for n in range(10):
    print(translator.train_step([input_texts3, target_texts3]))
print()

```

```

{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=6.247205>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=6.1790805>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=6.041966>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=5.6836166>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=4.7087393>}

```

```
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=4.7425594>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=3.9036467>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=3.8094451>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=3.8800523>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=3.7265575>}
```

```
CPU times: user 21min 1s, sys: 18.1 s, total: 21min 19s
Wall time: 11min 6s
```

```
@tf.function(input_signature=[tf.TensorSpec(dtype=tf.string, shape=[None]),
                             tf.TensorSpec(dtype=tf.string, shape=[None])])
def _tf_train_step(self, inputs):
    return self._train_step(inputs)
```

```
TrainTranslator._tf_train_step = _tf_train_step
```

```
translator.use_tf_function = True
```

```
translator.train_step([example_input_batch, example_target_batch])
```

```
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=6.270193>}
```

```
%%time
for n in range(10):
    print(translator.train_step([example_input_batch, example_target_batch]))
print()
```


```
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=6.202393>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=6.0730004>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=5.7399745>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=4.8307624>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=4.1018925>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=3.4566236>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=3.3494666>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=3.2292607>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=2.9973493>}
{'batch_loss': <tf.Tensor: shape=(), dtype=float32, numpy=2.7920635>}
```

```
CPU times: user 31.1 s, sys: 536 ms, total: 31.6 s
Wall time: 16.4 s
```

```
losses = []
for n in range(100):
    print('.', end='')
    logs = translator.train_step([example_input_batch, example_target_batch])
    losses.append(logs['batch_loss'].numpy())

print()
plt.plot(losses)
```

.....  
 [<matplotlib.lines.Line2D at 0x7f5128ab7510>]



```
train_translator = TrainTranslator(
    embedding_dim, units,
    input_text_processor=input_text_processor,
    output_text_processor=output_text_processor)
```

```
# Configure the loss and optimizer
train_translator.compile(
    optimizer=tf.optimizers.Adam(),
    loss=MaskedLoss(),
    metrics=["accuracy"]
)
```

```
class BatchLogs(tf.keras.callbacks.Callback):
    def __init__(self, key):
        self.key = key
        self.logs = []
```

```
    def on_train_batch_end(self, n, logs):
        self.logs.append(logs[self.key])
```

```
batch_loss = BatchLogs('batch_loss')
```

```
train_translator.fit(dataset, epochs=50,
                     callbacks=[batch_loss])
```

```
Epoch 1/50
47/47 [=====] - 86s 2s/step - batch_loss: 3.5051
Epoch 2/50
47/47 [=====] - 80s 2s/step - batch_loss: 2.5841
Epoch 3/50
47/47 [=====] - 79s 2s/step - batch_loss: 2.1906
Epoch 4/50
47/47 [=====] - 80s 2s/step - batch_loss: 1.9036
Epoch 5/50
```



```

47/47 [=====] - 79s 2s/step - batch_loss: 1.6102
Epoch 6/50
47/47 [=====] - 80s 2s/step - batch_loss: 1.3565
Epoch 7/50
47/47 [=====] - 81s 2s/step - batch_loss: 1.1152
Epoch 8/50
47/47 [=====] - 80s 2s/step - batch_loss: 0.9087
Epoch 9/50
47/47 [=====] - 80s 2s/step - batch_loss: 0.7269
Epoch 10/50
47/47 [=====] - 82s 2s/step - batch_loss: 0.5924
Epoch 11/50
47/47 [=====] - 80s 2s/step - batch_loss: 0.4673
Epoch 12/50
47/47 [=====] - 81s 2s/step - batch_loss: 0.3870
Epoch 13/50
47/47 [=====] - 80s 2s/step - batch_loss: 0.3124
Epoch 14/50
47/47 [=====] - 80s 2s/step - batch_loss: 0.2761
Epoch 15/50
47/47 [=====] - 80s 2s/step - batch_loss: 0.2409
Epoch 16/50
47/47 [=====] - 80s 2s/step - batch_loss: 0.2234
Epoch 17/50
47/47 [=====] - 80s 2s/step - batch_loss: 0.2124
Epoch 18/50
47/47 [=====] - 79s 2s/step - batch_loss: 0.2003
Epoch 19/50
47/47 [=====] - 79s 2s/step - batch_loss: 0.1925
Epoch 20/50
47/47 [=====] - 80s 2s/step - batch_loss: 0.1865
Epoch 21/50
47/47 [=====] - 79s 2s/step - batch_loss: 0.1803
Epoch 22/50
47/47 [=====] - 80s 2s/step - batch_loss: 0.1763
Epoch 23/50
47/47 [=====] - 79s 2s/step - batch_loss: 0.1773
Epoch 24/50
47/47 [=====] - 80s 2s/step - batch_loss: 0.1701
Epoch 25/50
47/47 [=====] - 80s 2s/step - batch_loss: 0.1742
Epoch 26/50
47/47 [=====] - 79s 2s/step - batch_loss: 0.1670
Epoch 27/50
47/47 [=====] - 79s 2s/step - batch_loss: 0.1666
Epoch 28/50
47/47 [=====] - 80s 2s/step - batch_loss: 0.1615
Epoch 29/50
47/47 [=====] - 80s 2s/step - batch_loss: 0.1602
Epoch 30/50


```

```

plt.plot(batch_loss.logs)
plt.ylim([0, 3])
plt.xlabel('Batch #')
plt.ylabel('CE/token')

```

Text(0, 0.5, 'CE/token')



The loss decrease very fast but then does not stabilize. Would have to cancel the noise. However this is a goog model with the lowest lose at .1413. The lowest out of the three models.