

```

import os
import tensorflow as tf
# Load compressed models from tensorflow_hub
os.environ['TFHUB_MODEL_LOAD_FORMAT'] = 'COMPRESSED'
import IPython.display as display
import matplotlib.pyplot as plt
import matplotlib as mpl
mpl.rcParams['figure.figsize'] = (12,12)
mpl.rcParams['axes.grid'] = False
import numpy as np
import PIL.Image
import time
import functools

```

```

def tensor_to_image(tensor):
    tensor = tensor*255
    tensor = np.array(tensor, dtype=np.uint8)
    if np.ndim(tensor)>3:
        assert tensor.shape[0] == 1
        tensor = tensor[0]
    return PIL.Image.fromarray(tensor)

```

```

from google.colab import drive
drive.mount('/content/drive')

```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mou



```

csueb_path='./content/students-walking-on-campus.png'
painted_path='./content/horse_painting.jpg'

```

```

def load_img(path_to_img):
    max_dim = 512
    img = tf.io.read_file(path_to_img)
    img = tf.image.decode_image(img, channels=3)
    img = tf.image.convert_image_dtype(img, tf.float32)
    shape = tf.cast(tf.shape(img)[: -1], tf.float32)
    long_dim = max(shape)
    scale = max_dim / long_dim
    new_shape = tf.cast(shape * scale, tf.int32)
    img = tf.image.resize(img, new_shape)
    img = img[tf.newaxis, :]
    return img

```

```

def imshow(image, title=None):
    if len(image.shape) > 3:
        image = tf.squeeze(image, axis=0)
    plt.imshow(image)



```

```

    if title:
        plt.title(title)

content_image = load_img(csueb_path)
style_image = load_img(painted_path)
plt.subplot(1, 2, 1)
imshow(content_image, 'CSUEB Image')
plt.subplot(1, 2, 2)
imshow(style_image, 'Painted Image')

```



```

import tensorflow_hub as hub
hub_model = hub.load('https://tfhub.dev/google/magenta/arbitrary-image-stylization-v1-256/2')
stylized_image = hub_model(tf.constant(content_image), tf.constant(style_image))[0]
tensor_to_image(stylized_image)
x = tf.keras.applications.vgg19.preprocess_input(content_image*255)
x = tf.image.resize(x, (224, 224))
vgg = tf.keras.applications.VGG19(include_top=True, weights='imagenet')
prediction_probabilities = vgg(x)
prediction_probabilities.shape
predicted_top_5 = tf.keras.applications.vgg19.decode_predictions(prediction_probabilities.numpy()
[(class_name, prob) for (number, class_name, prob) in predicted_top_5]
print()

```

Downloading data from <https://storage.googleapis.com/tensorflow/keras-applications/vgg1574717952/574710816> [=====] - 5s 0us/step

```
574726144/574710816 [=====] - 5s 0us/step
Downloading data from https://storage.googleapis.com/download.tensorflow.org/data/image40960/35363 [=====] - 0s 0us/step
49152/35363 [=====] - 0s 0us/step
```



```
for layer in vgg.layers:
    print(layer.name)
```

```
input_1
block1_conv1
block1_conv2
block1_pool
block2_conv1
block2_conv2
block2_pool
block3_conv1
block3_conv2
block3_conv3
block3_conv4
block3_pool
block4_conv1
block4_conv2
block4_conv3
block4_conv4
block4_pool
block5_conv1
block5_conv2
block5_conv3
block5_conv4
block5_pool
flatten
fc1
fc2
predictions
```

```
content_layers = ['block5_conv2']
style_layers = ['block1_conv1',
                'block2_conv1',
                'block3_conv1',
                'block4_conv1',
                'block5_conv1']
num_content_layers = len(content_layers)
num_style_layers = len(style_layers)
```

```
def vgg_layers(layer_names):
    """ Creates a vgg model that returns a list of intermediate output values. """
    # Load our model. Load pretrained VGG, trained on imagenet data
    vgg = tf.keras.applications.VGG19(include_top=False, weights='imagenet')
    vgg.trainable = False
```

```

outputs = [vgg.get_layer(name).output for name in layer_names]

model = tf.keras.Model([vgg.input], outputs)
return model

style_extractor = vgg_layers(style_layers)
style_outputs = style_extractor(style_image*255)

#Look at the statistics of each layer's output
for name, output in zip(style_layers, style_outputs):
    print(name)
    print("shape:", output.numpy().shape)
    print("min:", output.numpy().min())
    print("max:", output.numpy().max())
    print("mean:", output.numpy().mean())
    print()

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg1
80142336/80134624 [=====] - 1s 0us/step
80150528/80134624 [=====] - 1s 0us/step
block1_conv1
  shape: (1, 511, 408, 64)
  min: 0.0
  max: 875.37885
  mean: 35.754337

block2_conv1
  shape: (1, 255, 204, 128)
  min: 0.0
  max: 4549.171
  mean: 197.31772

block3_conv1
  shape: (1, 127, 102, 256)
  min: 0.0
  max: 11595.237
  mean: 185.18872

block4_conv1
  shape: (1, 63, 51, 512)
  min: 0.0
  max: 24262.912
  mean: 680.86816

block5_conv1
  shape: (1, 31, 25, 512)
  min: 0.0
  max: 2637.521
  mean: 45.93024

```

```

def gram_matrix(input_tensor):
    result = tf.linalg.einsum('biic biid->bd', input_tensor, input_tensor)

```

```

result = tf.nn.conv2d(input_image, filters, input_tensor, input_tensor,
input_shape = tf.shape(input_tensor)
num_locations = tf.cast(input_shape[1]*input_shape[2], tf.float32)
return result/(num_locations)

class StyleContentModel(tf.keras.models.Model):
    def __init__(self, style_layers, content_layers):
        super(StyleContentModel, self).__init__()
        self.vgg = vgg_layers(style_layers + content_layers)
        self.style_layers = style_layers
        self.content_layers = content_layers
        self.num_style_layers = len(style_layers)
        self.vgg.trainable = False

    def call(self, inputs):
        "Expects float input in [0,1]"
        inputs = inputs*255.0
        preprocessed_input = tf.keras.applications.vgg19.preprocess_input(inputs)
        outputs = self.vgg(preprocessed_input)
        style_outputs, content_outputs = (outputs[:self.num_style_layers],
                                          outputs[self.num_style_layers:])

        style_outputs = [gram_matrix(style_output)
                        for style_output in style_outputs]

        content_dict = {content_name: value
                        for content_name, value
                        in zip(self.content_layers, content_outputs)}

        style_dict = {style_name: value
                     for style_name, value
                     in zip(self.style_layers, style_outputs)}

        return {'content': content_dict, 'style': style_dict}

extractor = StyleContentModel(style_layers, content_layers)

results = extractor(tf.constant(content_image))

print('Styles:')
for name, output in sorted(results['style'].items()):
    print(".", name)
    print("....shape:", output.numpy().shape)
    print("....min:", output.numpy().min())
    print("....max:", output.numpy().max())
    print("....mean:", output.numpy().mean())
    print()

print("Contents:")
for name, output in sorted(results['content'].items()):
    print(".", name)

```

```

..print("....shape:", output.numpy().shape)
..print("....min:", output.numpy().min())
..print("....max:", output.numpy().max())
..print("....mean:", output.numpy().mean())

```

Styles:

```

block1_conv1
  shape: (1, 64, 64)
  min: 0.2309468
  max: 31140.148
  mean: 984.14825

```

```

block2_conv1
  shape: (1, 128, 128)
  min: 0.0
  max: 239140.47
  mean: 31675.844

```

```

block3_conv1
  shape: (1, 256, 256)
  min: 0.0
  max: 948065.1
  mean: 40235.055

```

```

block4_conv1
  shape: (1, 512, 512)
  min: 0.0
  max: 8885830.0
  mean: 542932.6

```

```

block5_conv1
  shape: (1, 512, 512)
  min: 0.0
  max: 476421.28
  mean: 5435.8745

```

Contents:

```

block5_conv2
  shape: (1, 16, 31, 512)
  min: 0.0
  max: 3961.6167
  mean: 29.52333

```

```

style_targets = extractor(style_image)['style']
content_targets = extractor(content_image)['content']

```

```

image = tf.Variable(content_image)

```

```

def clip_0_1(image):
    return tf.clip_by_value(image, clip_value_min=0.0, clip_value_max=1.0)

```

```

opt = tf.optimizers.Adam(learning_rate=0.02, beta_1=0.99, epsilon=1e-1)

```

```
style_weight=1e-2
content_weight=1e4
```

```
def style_content_loss(outputs):
    ...style_outputs.=.outputs['style']
    ...content_outputs.=.outputs['content']
    ...style_loss.=.tf.add_n([tf.reduce_mean((style_outputs[name]-style_targets[name])**2)
    .....for name in style_outputs.keys()])
    ...style_loss*=.style_weight./num_style_layers

    ...content_loss.=.tf.add_n([tf.reduce_mean((content_outputs[name]-content_targets[name])**2)
    .....for name in content_outputs.keys()])
    ...content_loss*=.content_weight./num_content_layers
    ...loss.=.style_loss+.content_loss
    ...return loss
```

```
@tf.function()
def train_step(image):
    ..with tf.GradientTape().as_tape:
    ...outputs.=.extractor(image)
    ...loss.=.style_content_loss(outputs)

    ..grad.=.tape.gradient(loss,.image)
    ..opt.apply_gradients([(grad,.image)])
    ..image.assign(clip_0_1(image))
```

```
train_step(image)
train_step(image)
train_step(image)
tensor_to_image(image)
```



```
import time
start.=.time.time()
```



```
start = time.time()
```

```
epochs = 10
steps_per_epoch = 100

step = 0
for n in range(epochs):
    for m in range(steps_per_epoch):
        step += 1
        train_step(image)
        print(".", end='', flush=True)
    display.clear_output(wait=True)
    display.display(tensor_to_image(image))
    print("Train step: {}".format(step))

end = time.time()
print("Total time: {:.1f}".format(end-start))
```




Train step: 1000  
Total time: 3972.1

## New Picture

```
csueb_path = '/content/Cal state east bay.jpg'
painted_path = '/content/painting_1.jpg'
```

```
content_image = load_img(csueb_path)
style_image = load_img(painted_path)
plt.subplot(1, 2, 1)
imshow(content_image, 'CSUEB Image 2')
plt.subplot(1, 2, 2)
imshow(style_image, 'Painted Image 2')
```





```
import tensorflow_hub as hub
hub_model = hub.load('https://tfhub.dev/google/magenta/arbitrary-image-stylization-v1-256/2')
stylized_image = hub_model(tf.constant(content_image), tf.constant(style_image))[0]
tensor_to_image(stylized_image)
x = tf.keras.applications.vgg19.preprocess_input(content_image*255)
x = tf.image.resize(x, (224, 224))
vgg = tf.keras.applications.VGG19(include_top=True, weights='imagenet')
prediction_probabilities = vgg(x)
prediction_probabilities.shape
predicted_top_5 = tf.keras.applications.vgg19.decode_predictions(prediction_probabilities.numpy()[0], top=5)
[(class_name, prob) for (number, class_name, prob) in predicted_top_5]
print()
```

```
for layer in vgg.layers:
    print(layer.name)
```

```
input_4
block1_conv1
block1_conv2
block1_pool
block2_conv1
block2_conv2
block2_pool
block3_conv1
block3_conv2
block3_conv3
block3_conv4
block3_pool
block4_conv1
block4_conv2
block4_conv3
block4_conv4
block4_pool
block5_conv1
block5_conv2
block5_conv3
block5_conv4
```

```
block5_pool  
flatten  
fc1  
fc2  
predictions
```

```
content_layers = ['block5_conv2']  
style_layers = ['block1_conv1',  
'block2_conv1',  
'block3_conv1',  
'block4_conv1',  
'block5_conv1']  
num_content_layers = len(content_layers)  
num_style_layers = len(style_layers)
```

```
style_extractor = vgg_layers(style_layers)  
style_outputs = style_extractor(style_image*255)
```

```
#Look at the statistics of each layer's output  
for name, output in zip(style_layers, style_outputs):  
    print(name)  
    print("  shape: ", output.numpy().shape)  
    print("  min: ", output.numpy().min())  
    print("  max: ", output.numpy().max())  
    print("  mean: ", output.numpy().mean())  
    print()
```

```
block1_conv1  
  shape: (1, 288, 512, 64)  
  min:  0.0  
  max:  824.818  
  mean:  34.34097
```

```
block2_conv1  
  shape: (1, 144, 256, 128)  
  min:  0.0  
  max:  4492.1304  
  mean:  221.71385
```

```
block3_conv1  
  shape: (1, 72, 128, 256)  
  min:  0.0  
  max:  9234.559  
  mean:  234.34587
```

```
block4_conv1  
  shape: (1, 36, 64, 512)  
  min:  0.0  
  max:  25270.893  
  mean:  850.7247
```

```
block5_conv1
```

```
shape: (1, 18, 32, 512)
min: 0.0
max: 3231.2678
mean: 58.030205
```

```
extractor = StyleContentModel(style_layers, content_layers)
```

```
results = extractor(tf.constant(content_image))
```

```
print('Styles:')
for name, output in sorted(results['style'].items()):
    print(" ", name)
    print("    shape: ", output.numpy().shape)
    print("    min: ", output.numpy().min())
    print("    max: ", output.numpy().max())
    print("    mean: ", output.numpy().mean())
    print()
```

```
print("Contents:")
for name, output in sorted(results['content'].items()):
    print(" ", name)
    print("    shape: ", output.numpy().shape)
    print("    min: ", output.numpy().min())
    print("    max: ", output.numpy().max())
    print("    mean: ", output.numpy().mean())
```

```
Styles:
  block1_conv1
    shape: (1, 64, 64)
    min: 0.10935415
    max: 36434.26
    mean: 490.3711

  block2_conv1
    shape: (1, 128, 128)
    min: 0.0
    max: 109473.48
    mean: 17873.855

  block3_conv1
    shape: (1, 256, 256)
    min: 0.0
    max: 518084.62
    mean: 16948.277

  block4_conv1
    shape: (1, 512, 512)
    min: 0.0
    max: 4669337.0
    mean: 242333.19

  block5_conv1
```

```

shape: (1, 512, 512)
min: 0.0
max: 112971.82
mean: 1811.5908

```

Contents:

```

block5_conv2
shape: (1, 12, 31, 512)
min: 0.0
max: 1096.959
mean: 16.499222

```

```

style_targets = extractor(style_image)['style']
content_targets = extractor(content_image)['content']

```

```

image2 = tf.Variable(content_image)

```

```

opt = tf.optimizers.Adam(learning_rate=0.02, beta_1=0.99, epsilon=1e-1)
style_weight=1e-2
content_weight=1e4

```

```

def style_content_loss(outputs):
    style_outputs = outputs['style']
    content_outputs = outputs['content']
    style_loss = tf.add_n([tf.reduce_mean((style_outputs[name]-style_targets[name])**2)
                           for name in style_outputs.keys()])
    style_loss *= style_weight / num_style_layers

    content_loss = tf.add_n([tf.reduce_mean((content_outputs[name]-content_targets[name])**2)
                             for name in content_outputs.keys()])
    content_loss *= content_weight / num_content_layers
    loss = style_loss + content_loss
    return loss

```

```

@tf.function()
def train_step(image):
    with tf.GradientTape() as tape:
        outputs = extractor(image)
        loss = style_content_loss(outputs)


    grad = tape.gradient(loss, image)
    opt.apply_gradients([(grad, image)])
    image.assign(clip_0_1(image))

```

```

train_step(image2)
train_step(image2)
train_step(image2)
tensor_to_image(image2)

```




```
import time
start = time.time()

epochs = 10
steps_per_epoch = 100

step = 0
for n in range(epochs):
    for m in range(steps_per_epoch):
        step += 1
        train_step(image2)
        print(".", end='', flush=True)
        display.clear_output(wait=True)
        display.display(tensor_to_image(image2))
        print("Train step: {}".format(step))

end = time.time()
print("Total time: {:.1f}".format(end-start))
```



Train step: 1000  
Total time: 3192.3

---

✓ 53m 12s    completed at 9:48 PM ● ✕