

DNN Lab

Objectives

- Understand basic DNN model building process using Keras
- Analyze model performance and capacity vs generalization tradeoff
- Modify models to reduce overfitting and improve performance

Exercises

- Build a DNN model for slump Test Problem
- Start with a model consisting of one hidden layer with 7 neurons
- Analyze results and explore improvements to model in terms of capacity, regularization

Step 1: Import Libraries

```
In [112]: %tensorflow_version 2.x
          from numpy.random import seed
          seed(2)
          import tensorflow as tf
          from tensorflow import keras
          from IPython import display
          from matplotlib import cm
          from matplotlib import gridspec
          from matplotlib import pyplot as plt
          import numpy as np
          import pandas as pd
          import os
          import datetime
          from tensorflow.python.data import Dataset
          from sklearn import preprocessing
          from sklearn.preprocessing import StandardScaler, StandardScaler
          from sklearn.model_selection import train_test_split

          print(tf.__version__)
```

2.6.0

Step 2: Import Data

```
In [113]: pd.options.display.max_rows = 10
pd.options.display.float_format = '{:.1f}'.format

hcv_data = pd.read_csv("hcvdat0.csv")

hcv_data = hcv_data.reindex(
    np.random.permutation(hcv_data.index))
```

```
In [114]: hcv_data.shape[0]
```

```
Out[114]: 615
```

```
In [115]: #removing the redundant index column
hcv_data.drop('Unnamed: 0', axis=1, inplace=True)
print(hcv_data)
```

	Category	Age	Sex	ALB	ALP	...	CHE	CHOL	CREA	GGT
PROT										
469	0=Blood Donor	52	f	51.5	81.8	...	6.7	5.9	88.0	16.3
82.2										
592	3=Cirrhosis	47	m	42.0	nan	...	6.3	5.5	58.0	201.0
79.0										
265	0=Blood Donor	58	m	41.3	58.9	...	8.2	5.7	60.0	10.8
70.1										
84	0=Blood Donor	39	m	43.9	90.1	...	9.9	4.6	98.0	99.3
66.2										
109	0=Blood Donor	42	m	44.1	46.8	...	10.8	6.3	95.0	19.7
73.0										
..
...										
534	0s=suspect Blood Donor	48	m	24.9	116.9	...	3.4	5.2	29.0	83.0
47.8										
584	2=Fibrosis	75	f	36.0	nan	...	6.7	nan	57.0	177.0
72.0										
493	0=Blood Donor	56	f	34.7	90.3	...	8.1	5.5	67.0	9.0
69.4										
527	0=Blood Donor	63	f	27.8	85.7	...	6.1	4.0	63.0	46.0
56.9										
168	0=Blood Donor	47	m	48.3	59.3	...	11.1	5.6	88.0	91.5
73.0										

```
[615 rows x 13 columns]
```

```
In [116]: #2.3 Minimum accuracy the model needs to be
naive_app_min= hcv_data.Category.value_counts().max()/len(hcv_data)
naive_app_min
```

```
Out[116]: 0.8666666666666667
```

```
In [117]: #changing Sex column to 0 and 1s
hcv_data= pd.get_dummies(hcv_data, columns=['Sex'], drop_first=True)
#changing predictor variable into dummy vars
hcv_data= pd.get_dummies(hcv_data, columns=['Category'])
print(hcv_data)
```

	Age	ALB	...	Category_2=Fibrosis	Category_3=Cirrhosis
469	52	51.5	...	0	0
592	47	42.0	...	0	1
265	58	41.3	...	0	0
84	39	43.9	...	0	0
109	42	44.1	...	0	0
..
534	48	24.9	...	0	0
584	75	36.0	...	1	0
493	56	34.7	...	0	0
527	63	27.8	...	0	0
168	47	48.3	...	0	0

[615 rows x 17 columns]

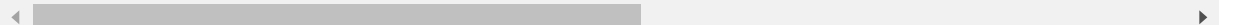
```
In [118]: hcv_data = hcv_data[hcv_data.columns[:-1]]
```

Step 3: Preprocess

```
In [119]: #removed any NaN values and replaced it with the column mean
for cols in hcv_data.columns[hcv_data.isnull().any()]:
    hcv_data[cols].replace(np.NaN, hcv_data[cols].mean(),inplace=True)

hcv_data[hcv_data.isnull().any(axis=1)]
```

Out[119]:

Category_3=Cirrhosis	Category_2=Fibrosis	Category_1=Hepatitis	Category_0s=suspect Blood Donor	Category_0=Healthy
				

```
In [120]: #checking to see if any NaN values are present
len(hcv_data[hcv_data.isnull().any(axis=1)])
```

Out[120]: 0

Train/Validation Split

```
In [121]: #Creating a training and validation dataset with a 80/20 split
X_train,X_test, y_train, y_test = train_test_split(hcv_data.iloc[:,5:],hcv_data.
a.iloc[:,5], test_size=0.2, random_state=1)
```

In [122]: *#2.2 Print first few rows of training data*
 X_train.head()

Out[122]:

	Sex_m	PROT	GGT	CREA	CHOL	CHE	BIL	AST	ALT	ALP	ALB	Age
246	1	72.2	30.2	80.0	6.3	7.5	4.5	21.0	36.9	87.1	46.2	55
344	0	72.8	12.4	64.0	5.1	10.0	3.0	22.0	19.2	62.6	43.4	35
75	1	70.1	17.3	67.0	4.1	6.9	16.7	35.1	47.4	69.4	44.7	38
216	1	67.4	87.8	77.0	6.1	8.9	7.8	23.7	37.0	82.2	82.2	52
444	0	78.2	14.5	69.0	6.8	9.1	5.8	15.9	14.3	45.9	45.4	49

In [123]: *#normalizing training dataset*
 scaler = StandardScaler()

 scaledf = scaler.fit_transform(X_train.iloc[:,1:])
 X_train.iloc[:,1:] = pd.DataFrame(scaledf, index=X_train.iloc[:,1:].index, columns=X_train.iloc[:,1:].columns)

#print(X_train)
#normalizing validation dataset
 vscaled = scaler.transform(X_test.iloc[:,1:].values)
 X_test.iloc[:,1:] = pd.DataFrame(vscaled, index=X_test.iloc[:,1:].index, columns=X_test.iloc[:,1:].columns)
#print(X_test)

Step 4: Build Model

https://www.tensorflow.org/api_docs/python/tf/keras/Model
 (https://www.tensorflow.org/api_docs/python/tf/keras/Model)

https://www.tensorflow.org/api_docs/python/tf/keras/layers/Dense
 (https://www.tensorflow.org/api_docs/python/tf/keras/layers/Dense)

<https://keras.io/optimizers/> (<https://keras.io/optimizers/>)

Build Model

```

In [124]: l2_model = keras.Sequential([
            keras.layers.Dense(24, kernel_regularizer=keras.regularizers.l2(0.001), ac
            tivation=tf.nn.relu,
                                input_shape=(X_train.shape[1],)),
            keras.layers.Dropout(0.25),
            keras.layers.Dense(5, kernel_regularizer=keras.regularizers.l2(0.001), act
            ivation=tf.nn.relu),
            keras.layers.Dropout(0.25),
            #keras.layers.Dense(3, kernel_regularizer=keras.regularizers.l2(0.01), act
            ivation=tf.nn.relu),
            #keras.layers.Dropout(0.25),
            #keras.layers.Dense(6, kernel_regularizer=keras.regularizers.l2(0.01), act
            ivation=tf.nn.softmax),
            #keras.layers.Dropout(0.50),
            #keras.layers.Dense(4, kernel_regularizer=keras.regularizers.l2(0.01), act
            ivation=tf.nn.softmax),
            #keras.layers.Dropout(0.50),
            keras.layers.Dense(5, activation=tf.nn.softmax)
        ])

l2_model.compile(loss=tf.keras.losses.CategoricalCrossentropy(label_smoothing=
0.1),
                 optimizer='sgd',
                 metrics=[tf.keras.metrics.CategoricalAccuracy()])

```

Fit Model

```

In [125]: logdir = os.path.join("logs", datetime.datetime.now().strftime("%Y%m%d-%H%M%S"
))
tensorboard_callback = tf.keras.callbacks.TensorBoard(logdir, histogram_freq=
1)

```

In [126]: `l2_model.summary()`

Model: "sequential_7"

Layer (type)	Output Shape	Param #
=====		
dense_22 (Dense)	(None, 24)	312

dropout_15 (Dropout)	(None, 24)	0

dense_23 (Dense)	(None, 5)	125

dropout_16 (Dropout)	(None, 5)	0

dense_24 (Dense)	(None, 5)	30
=====		
Total params: 467		
Trainable params: 467		
Non-trainable params: 0		

```
In [ ]: class PrintDot(keras.callbacks.Callback):
        def on_epoch_end(self, epoch, logs):
            if epoch % 10 == 0:
                print('')
                print(logs)

EPOCHS = 500
tf.random.set_seed(1)

# Store training stats
l2_history = l2_model.fit(X_train, y_train, epochs=EPOCHS,
                          validation_data= (X_test, y_test), verbose=2) #PRINTDOTS NO
T WORKING RIGHT NOW
#try 0 and 1, verboe 1 gives data as its running
```

Step 5: Plot Results

In [128]: *#2.4 Chart of training and validation error and accuracy (or appropriate metric based on problem)*

```
import matplotlib.pyplot as plt

def plot_history(histories, key='loss'):
    plt.figure(figsize=(16,10))
    for name, history in histories:
        val = plt.plot(l2_history.epoch, l2_history.history['val_'+key],
                      '--', label=name.title()+' Val')
        plt.plot(l2_history.epoch, l2_history.history[key], color=val[0].get_color(),
                 label=name.title()+' Train')


    plt.xlabel('Epochs')
    plt.ylabel(key.replace('_', ' ').title())
    plt.legend()

    plt.xlim([0,max(l2_history.epoch)])
    plt.ylim([0,2])

plot_history([('Basic Model', l2_history)])

#Plot Multiple Model Results

#plot_history([('Plain', m1_history),('L1',model1)])
```



Predictions

```
In [138]: valpreds = np.round(l2_model.predict_on_batch(X_test),3)
print(valpreds[:5])
```


```
[[0.015 0.024 0.029 0.02  0.913]
 [0.016 0.025 0.029 0.021 0.909]
 [0.021 0.031 0.033 0.021 0.894]
 [0.018 0.028 0.031 0.02  0.903]
 [0.016 0.025 0.03  0.019 0.911]]
```

```
In [ ]: with pd.option_context('display.max_rows', None, 'display.max_columns', None):
print(y_test)
```

```
In [130]: # Plot Weights
nfw = l2_model.get_weights()[0][0]
y_pos = np.arange(len(nfw))

plt.bar(y_pos, nfw, align='center', alpha=0.5)
```

Out[130]: <BarContainer object of 24 artists>



```
In [133]: print(min(l2_history.history['val_loss']))

0.6051108837127686
```

```
In [135]: print(max(l2_history.history['val_categorical_accuracy']))

0.8943089246749878
```

```
In [ ]: 1
```


Goal: Predict whether the patient belongs in 1 of 5 categories: Blood Donor, Suspected Blood Donor, Hepatitis, Fibrosis, and Cirrhosis. Given the 5 categories we're trying to predict the a model >naive approach accuracy of 86.67%.

As you can see in the code above the categorical accuracy was 89.43%, which is slightly higher than the baseline accuracy of 86.67%.

The graph shows very little generalization issues since they are converging. Label Smoothing made it smoother then prior because the spikes were frequent. I also had to use less layers as the accuracy was pretty low and the generalization error was significant.

Changed the amount of nodes as a larger amount of nodes had a an accuracy below the baseline.